

Assignment Project Exam Help

<https://eduassistpro.github.io/> elon University

Add WeChat edu_assist_pro

How to Write Fast Code?

Fast Platforms



Good Techniques

Assignment Project Exam Help

- Multicore platforms
 - Manycore
 - Cloud plattform
- <https://eduassistpro.github.io/>

- Data structures

- Course Goals

Add WeChat edu_assist_pro

- To write fast code for your research/applications
 - 1. Feel comfortable hacking up a solution
 - 2. Leverage existing software building blocks
 - 3. Indicate which platform is the best one to use
 - 4. Reason about why a piece of existing code is slow
 - 5. Take care of potential performance bottlenecks

Course Staff

- **Instructor: Ian Lane**

- Preferred contact method: [Piazza](#)
- [Office Hours](#): Wed 4:30-5:30pm ET
- Email: ianl@mit.edu

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- **Teaching Assistants**

- Deepak Arumugam Sankara Subramanianh
- Yunfan Jiang (Silicon Valley)

Course Information

- **Lectures:**
 - Tuesday and Thursday 6:00pm-7:20pm ET
- **Office Hours:**
 - Instructor Office Hours: Wednesdays 4:30pm-5:30pm ET
 - TA Office Ho
- **Grading:** <https://eduassistpro.github.io/>
 - 10% - Homeworks
 - 30% - Mini-Projects
 - 30% - Term Project
 - 30% - Final Exam
- **Course Links:**
 - **Canvas:** <https://canvas.cmu.edu/courses/21510/pages/course-schedule>
 - **Piazza:** <https://piazza.com/class/kkmp02yc92h598>
 - **GradeScope:** <https://www.gradescope.com/courses/241050>

<https://canvas.cmu.edu/courses/21510/pages/course-schedule>

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

<https://piazza.com/class/kkmp02yc92h598>

Sign Up Link: <https://piazza.com/cmu/spring2021/18646>

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

<https://www.gradescope.com/courses/241050>

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Outline

- Why Fast Code Now?
- What does 100x speed mean?
Assignment Project Exam? Help
- Course structure
<https://eduassistpro.github.io/>
- Problem sets
Add WeChat edu_assist_pro

Why Fast Code?

Need is driven by the applications...

...NOT by the availability of the platforms

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

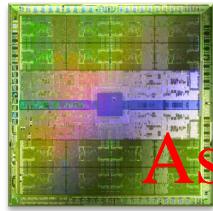
RMS Appears Everywhere

- **Recognition:** Interpretations of the world (with models)
- **Mining:** Understanding of the world (discover hidden patterns with model)
- **Synthesis:** Anticipate outcomes in the world (using the models to predict)

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

What is being valued today?



Hardware → Commoditized
Assignment Project Exam Help

```
int main(void) {  
    mykernel<<<1,1>>>();  
    printf("Hello World!  
    return 0;  
}
```

<https://eduassistpro.github.io/>



Add WeChat edu_assist_pro



Data



King!

Outline

- Why Fast Code Now?
 - What does 100x speedup mean? Assignment Project Exam Help
 - Course structure
 - Problem sets
- <https://eduassistpro.github.io/>
- Add WeChat edu_assist_pro

What does 100x Speedup Mean?

12 hours: 10x speedup → 1.2 hours

100x speedup → 7.2 minutes

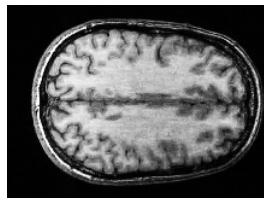
1000x speedup → 43.2 seconds

<https://eduassistpro.github.io/>

- Game changing
 - Overnight jobs becomes interactive



Speech Analytics



Medical Imaging

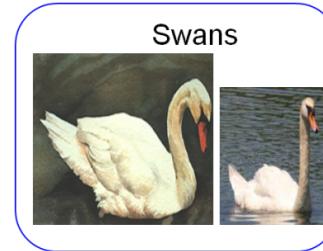


Image Recognition



Computational Finance

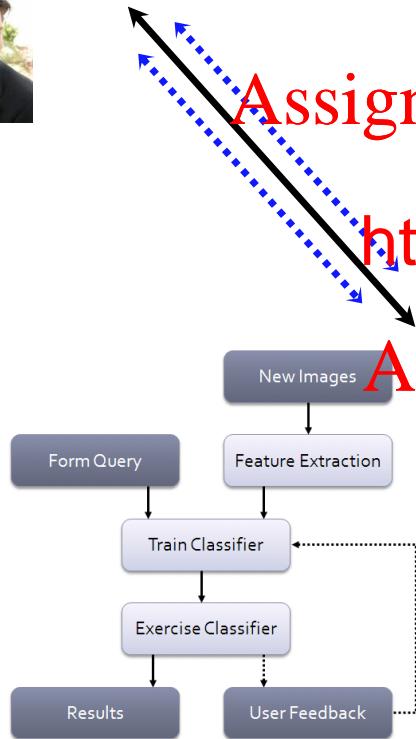
Fast, Robust Pediatric MRI

- **100X** faster reconstruction
- Higher-quality, faster MRI
- This image: 8 month-old patient with cancerous mass in
 - $256 \times 84 \times 154 \times$ <https://eduassistpro.github.io/>
 - Serial Recon: 1 hour
 - Parallel Recon: less than 1 minute
- Fast enough for clinical use
 - Software currently deployed at Lucile Packard Children's Hospital for clinical study of the reconstruction technique

Support-Vector Machines

- Algorithmic changes and parallel implementation lead to performance

Speed-up: Core 2 Duo vs G80



Assignment Project Exam Help

<https://eduassistpro.github.io/>

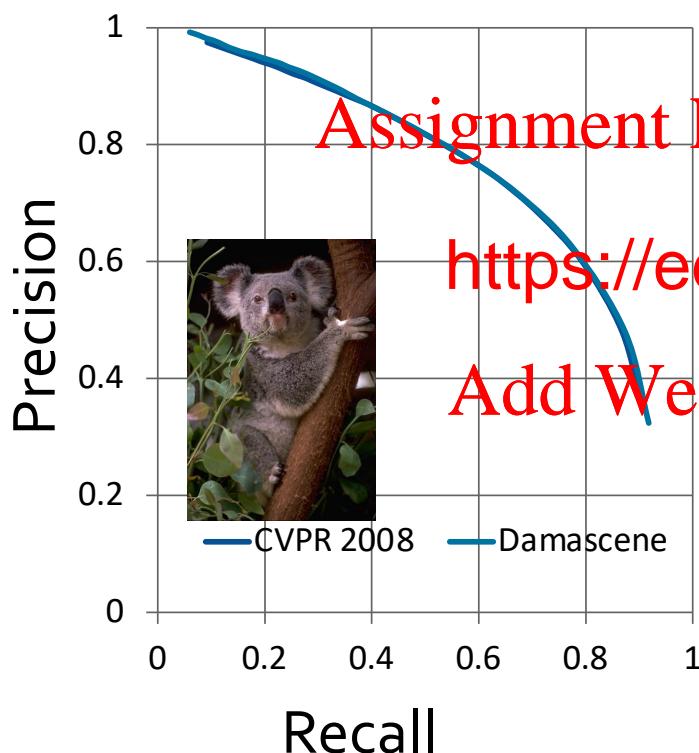
Add WeChat edu_assist_pro

| Procedure | Core 2 Duo | G80 | Our algorithm, 16-core Parallel GPU |
|-------------------------|------------|--------|-------------------------------------|
| SVM Training (geo-mean) | 38.5 s | 38.5 s | 0.38 s |
| SVM Classification | 41.42 s | 4.21 s | 0.38 s |

100X speed-up

896 downloads since release in 10/2008

Image Contours Detection



- We achieve equivalent accuracy on the Berkeley Segmentation Dataset

Comparing to human segmented
“ground truth”

or both

$H_t = 0.79$

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

records:

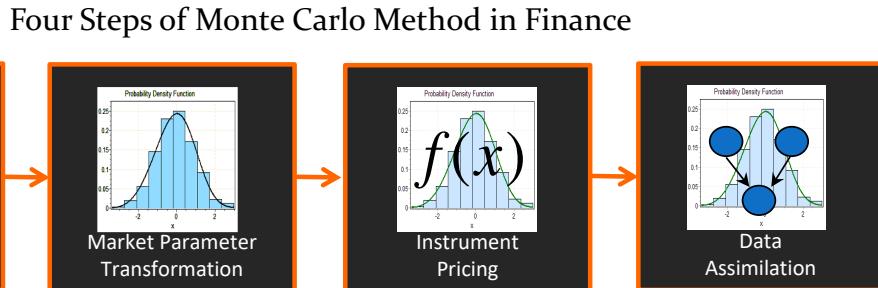
126x speedup

- 616 downloads since release in October 2009

Computational Finance

- Value-at-Risk Computation with Monte Carlo Method
- Summarizes a portfolio's vulnerabilities to market movements
- Important to alg derivative usage <https://eduassistpro.github.io/>
- Improved implementation to run 60x faster on a parallel microprocessor

Assignment Project Exam Help
Add WeChat [edu_assist_pro](#)



NY Times TIFF to PDF

- In 2007, the New York Times decided to make all the public domain articles from 1851-1922 available free of charge
 - Needed to convert from TIFF to PDF

Assignment Project Exam Help

Data set:

illion PDF
<https://eduassistpro.github.io/>

Compute Instances: 100 Amazon EC

Add WeChat edu_assist_pro

Time taken: 24 hours

Cost: \$240



<http://open.blogs.nytimes.com/2007/11/01/self-service-prorated-super-computing-fun/>

Assignment Project Exam Help

<https://eduassistpro.github.io/>

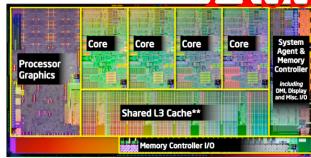
Add WeChat edu_assist_pro

Philosophy on Platforms

Provide theoretical background and hands-on practices...

...to innovate with multicore/manycore/cloud-based platforms.

Assignment Project Exam Help



Intel Sandy Bridge Multicore
Processor (Core i7-2600K)

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

NVIDIA Fermi Manycore
Processor - GTX580

Yahoo! Hadoop Cluster
~2000 nodes in one cluster

- ***Significant scientific advances*** will be empowered by multicore/manycore/cloud-based platforms over the next decade
- Knowledge of these new computing capabilities will give you an ***advantage*** over your peers in developing innovative techniques to solve challenging problems

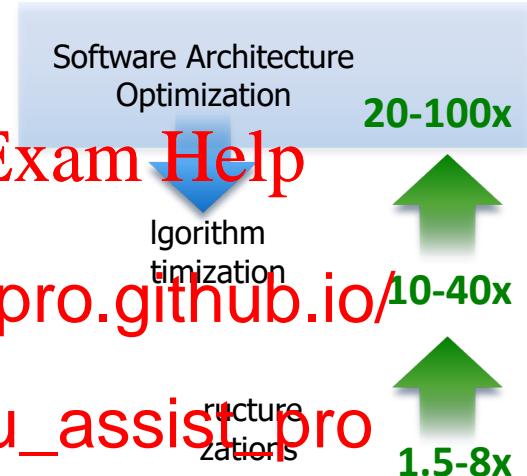
Philosophy on Techniques

- Efficient ***software architecture*** is the most important to designing fast code
 - Software design patterns
- Understanding the platform will help identify application performance bottlenecks
- ***Hands-on experience*** provide confidence for you to effectively use these technologies for your research and development needs

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Outline

- Why Fast Code Now?
- What does 100x speedup mean?
Assignment Project Exam? Help
- Course structure
<https://eduassistpro.github.io/>
- Problem sets
Add WeChat edu_assist_pro

<https://canvas.cmu.edu/courses/21510/quizzes/55032>

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Outline

- Why Fast Code Now?
- What does 100x speedup mean?
Assignment Project Exam? Help
- Course structure
<https://eduassistpro.github.io/>
- Problem sets
Add WeChat edu_assist_pro

Organization of This Course

- Lectures
 - Tuesdays and Thursdays
 - All **mini-project introductions** held on Tuesdays, **mini-project reviews** on Thursdays
 - Keep an eye on the course schedule for changes
- Canvas - <https://course-schedule>
 - Course content
 - Homework and Mini-Project Descriptions
- Piazza - <https://piazza.com/class/kkmp02yc92h598>
 - Course Announcements
 - Course Questions and Discussions
- Gradescope - <https://www.gradescope.com/courses/241050>
 - Assignment Submissions
 - Final Exam

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Organization of This Course

- During class:
 - Lectures and Q&A sessions
 - Outside class:
 - 3 homework assignments
 - 3 mini homeworks
 - 1 term project (focused on your own research)
 - Lots of coding ;)
- Assignment Project Exam Help**
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

Grading and Expectations

- GRADING

- 10% Homework assignments
- 30% Homework projects
- 30% Term Project
- 30% Final exam

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- EXPECTATIONS

- Attend majority of lectures
- Hand in all assignments and mini-projects
- Complete a term project
- Complete the final exam

Add WeChat edu_assist_pro

Software Reuse and Plagiarism

- **Question:** Can I learn to use code from the web?
- **Answer:** Yes! But only under certain circumstances...

| Plagiarism | Acceptable SW Reuse |
|--|---|
| Code Copying Unacknowledged | Code-level Reuse Refactored out to a separate file |
| Reverse Engineering Unacknowledged re-use of some code abstraction | Documentation in Documentation distinguished in documentation |
| Translation Unacknowledged re-use by translating functions from one language to another | Code Testing Third-party code must be tested to one's own requirements |
| Code Generation Unacknowledged assistance by using some code generators | |
| Reuse Without Testing Reuse third-party code without testing against one's own requirements | |

Gibson, "Software Reuse and Plagiarism: A Code of Practice", In Proceedings of ITiCSE'2009. pp.55-59

Structure of Lectures

- **Module 1: Background and Multicore Programming**
 - Hardware architectures, applications
 - Application design with OpenMP
- **Module 23: Man** <https://eduassistpro.github.io/>
 - Manycore architectures
 - Application design with CUDA
- **Module 3: Cluster Programming**
 - Distributed architectures
 - Application design with Hadoop

Assignment Project Exam Help

Add WeChat edu_assist_pro

Homework & Mini-Projects (40%)

- Three mini projects throughout the semester
 1. Multicore project with OpenMP
 2. Manycore project with CUDA
 3. Cloud project with Hadoop

Assignment Project Exam Help

- Projects are done at <https://eduassistpro.github.io/>
 - Programming assignments will start with template files
 - Accelerate them with techniques discussed in class
 - Will submit code, a project write-up and slides during class
-
- Homeworks are done in preparation for the Mini-Projects
 - Internalize concepts discussed in class
 - Set up the project environment
 - Discussion among peers is allowed and encouraged. However, tasks and write-ups must be completed individually

Homework & Mini-Project Schedule

- **Homeworks (Project Setup)** 10%
 - Module 1 – Multicore Thursday, Feburary 25th
 - Module 2 – Manycore Thursday, March 18th
 - Module 3 – , April 15th
- **Assignment Project Exam Help** Add WeChat `edu_assist_pro` 30%
 - Module 1 – Multicore Monday, March 8th
 - Module 2 – Manycore Monday, April 5th
 - Module 3 – Cluster Computing Monday, April 26th

Term Project (30%)

- **Term project**

- An application area of your choice
- Term project introduction on Tuesday next week (1/9)
- Projects are
- Does not ne

Assignment Project Exam Help

<https://eduassistpro.github.io/>
Project

Add WeChat edu_assist_pro

Term Project Schedule

- Term (Team) Project 30%
 - Project Proposal (10%) Tuesday, March 9th
 - Poster Presentations (40%) Tuesday, May 4th
 - Final Report May 14th
- Assignment Project Exam Help**
<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Exam (30%)

- Final Exam – TBD (Week of May 10th-14th)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Outline

- Why Fast Code Now?
- What does 100x speedup mean?
Assignment Project Exam? Help
- Course structure
<https://eduassistpro.github.io/>
- Problem sets
Add WeChat edu_assist_pro

Problem Solving for Fast Code

- Writing fast code is a process coherent with

“general problem solving behavior”

Assignment Project Exam Help

Newell and Simon, Human Problem Solving (1972), pp. 72-73

- The process of problem solving is:
<https://eduassistpro.github.io/>
 1. Understand the **current state**
 2. Observe the **internal representation**
 3. **Search** among alternatives
 4. Select from a set of **choices**

Add WeChat **edu_assist_pro**

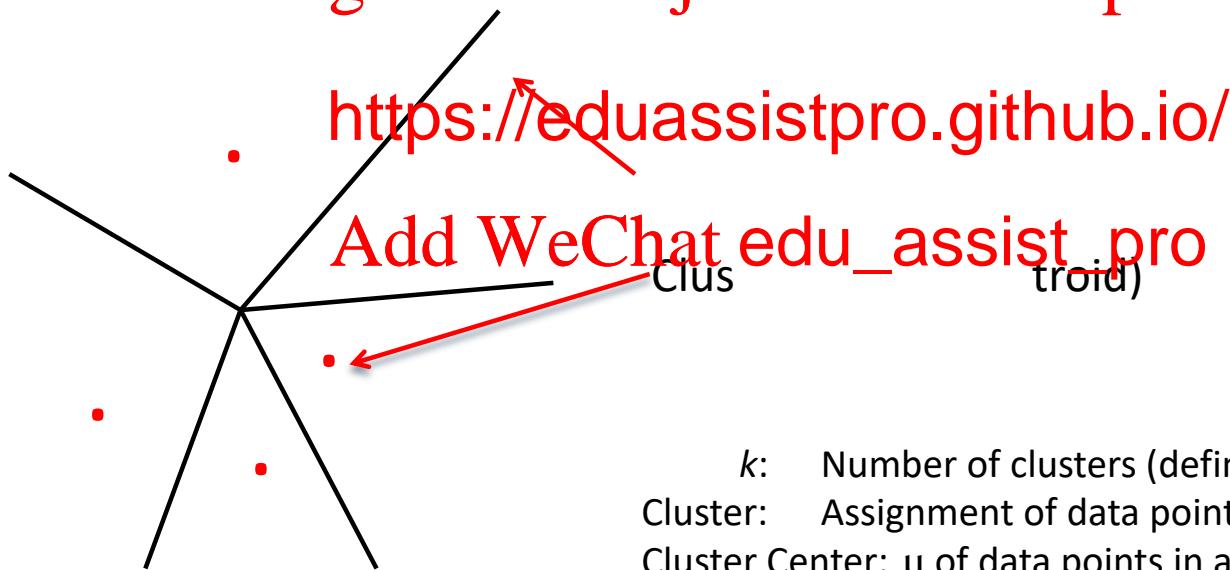
The k -means Problem

- Find k cluster centers that minimize the distance from each data point to a cluster center
- Important algorithm in machine learning:
 - Statistical data
 - Vector quantization
- NP-hard for arbitrary input
- **k -means algorithm** frequently finds a solution quickly
 - Add WeChat `edu_assist_pro`
- Issues:
 - Worst case running time is super-polynomial
 - Approximation can be arbitrarily bad

The k -means Problem

- Find k cluster centers that minimize the distance from each data point to a cluster center

Assignment Project Exam Help



k : Number of clusters (defined a-priori)
Cluster: Assignment of data points to a class
Cluster Center: μ of data points in a cluster

k-means Algorithm ("Lloyd's algorithm")

- Given an initial set of k means $\mathbf{m}_1^{(1)}, \dots, \mathbf{m}_k^{(1)}$
- Expectation Step: Assign each observation to the cluster with the closest mean

<https://eduassistpro.github.io/>

- Maximization Step: Calculate the centroid of the observations in the cluster.

$$\mathbf{m}_i^{(t+1)} = \frac{1}{|S_i^{(t)}|} \sum_{\mathbf{x}_j \in S_i^{(t)}} \mathbf{x}_j$$

- Iterate until convergence or stopping criteria met

The Algorithm

Example:

k=5

Distance metric=euclidean

Dimensions=2

1. Randomly select cluster Centers
2. Assign closest Center to each data point
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The Algorithm

Example:

k=5

Distance metric=euclidean

Dimensions=2

1. Randomly select cluster Centers
2. Assign closest Center to each data point
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The Algorithm

Example:

k=5

Distance metric=euclidean

Dimensions=2

1. Randomly select cluster Centers
2. Assign each data point to closest Center
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The Algorithm

Example:

k=5

Distance metric=euclidean

Dimensions=2

1. Randomly select cluster Centers
2. Assign closest Center to each data point
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The Algorithm

Example:

k=5

Distance metric=euclidean

Dimensions=2

1. Randomly select cluster Centers
2. Assign closest Center to each data point
3. Update Centers based on assignments from (2)
4. Re-iterate steps 2-3 until convergence or stopping criteria met

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The Phases

1. **Initialization:** Randomly select k cluster centers
 - Select k samples from data as initial centers [Forgy Partition]
2. **Expectation:** Assign each data point go closest center
 - Compare e
 - Distance M
3. **Maximization:** Update centers
 - For each cluster (k) compute mean (m_k) of all points assigned to that cluster
4. **Evaluate:** Re-iterate steps 2-3 until convergence or stopping criteria met
 - Percentage of data points re-assigned
 - Number of iterations (2-3)

Assignment Project Exam Help

<https://eduassistpro.github.io/>



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

A Fast Implementation of k -means

- Writing fast code is a process coherent with

“general problem solving behavior”

Assignment Project Exam Help

Newell and Simon, Human Problem Solving (1972), pp. 72-73

- The process of problem solving
<https://eduassistpro.github.io/>
 1. Understand the current state
 2. Observe the internal representation
 3. Search among alternatives
 4. Select from a set of choices

Add WeChat edu_assist_pro

A Fast Implementation of k -means

- Following the process of problem solving with k -means:

1. Understand the **current state**

- Running on a platform
- Using a spec
- Achieving a
- Meeting a specific criteria/requirement

Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

Assumption:

Starting from a functionally correct reference implementation

1: Observe the *current state* and
n requirements
to solve a problem

2. Observe the **internal representation**

3. **Search** among alternatives
4. Select from a set of **choices**

A Fast Implementation of k -means

- Following the process of problem solving with k -means:
 1. Understand the **current state**
 2. Observe the **internal representation**
 - Application
 - Identified <https://eduassistpro.github.io/>
 - Implementation concerns
 - Task considerations
 - Data representations
 - Concurrency opportunities
 3. **Search** among alternatives
 4. Select from a set of **choices**

A Fast Implementation of k -means

- Following the process of problem solving with k -means:
 1. Understand the **current state**
 2. Observe the **internal representation**
 - Application
 - Identified <https://eduassistpro.github.io/>
 - Implementation concerns
 - Task considerations
 - Data representations
 - Concurrency opportunities
 3. **Search among alternatives**
 4. Select from a set of **choices**

A Fast Implementation of k -means

- Following the process of problem solving with k -means:
 1. Understand the **current state**
 2. Observe the **internal representation**
 - Application
 - Identified <https://eduassistpro.github.io/>
 - Implementation concerns
 - Task considerations
 - Data representations
 - Concurrency opportunities
 3. **Search** among alternatives
 4. Select from a set of **choices**

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro