

Assignment Project Exam Help

<https://eduassistpro.github.io/>
Carnegie Mellon University

Add WeChat edu_assist_pro

Ian Lane

What we discussed last time:

Fast Platforms

Good Techniques

Assignment Project Exam Help

- Multicore p
- Manycore
- Cloud platf

ctures

<https://eduassistpro.github.io/>
Architecture

Add WeChat edu_assist_pro

- Highlighted the difference between mu core platforms
- Discussed the multicore and manycore platform intricacies
- Exposing concurrency in the k-means algorithm
- Exploiting parallelism by exploring mappings from application to platform

Example 7: Sequential Optimization

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main () {
    int i; double p
    step = 1.0/(dou
    #pragma omp par
    {te(x
        , local_sum)
    {
        int i, id = omp_get_thread_num()
        int nthrds = omp_get_num_threads()
        double hoop = nthrds*step;
        for (i=id, x=(i+0.5)*step;i< num_steps; i=i+nthrds) {
            x += hoop;
            local_sum += 4.0/(1.0+x*x);
        }
        #pragma omp critical
        pi += local_sum * step;
    }
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds

0.46 seconds

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Outline

- Multicores Shared-Memory Platforms
- OpenMP – Parallel Shared-Memory Multiprocessing
 - Overview
 - Example: numeric integration
 - Shared variable
 - **SPMD vs worksharing**
 - Data environment options
 - Advanced worksharing options
 - Synchronization options
- Sequential Optimizations
 - Cost Model
 - Cost Reduction

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

SPMD vs. Worksharing

- A parallel construct by itself creates an SPMD
 - “Single Program Multiple Data”
 - Programmer must explicitly specify what each thread must do differently
 - The division of work is hard-coded in the program
- Opportunity:
<https://eduassistpro.github.io/>
 - Many parallel regions are loops
- Question:
 - Can we make it easier to parallelize these loops?

OpenMP: Parallel For

- One of the worksharing constructs OpenMP provide is “**omp for**”

Sequential code

```
for(i=0;i<N;i++) { a[i] = a[i] + b[i];}
```

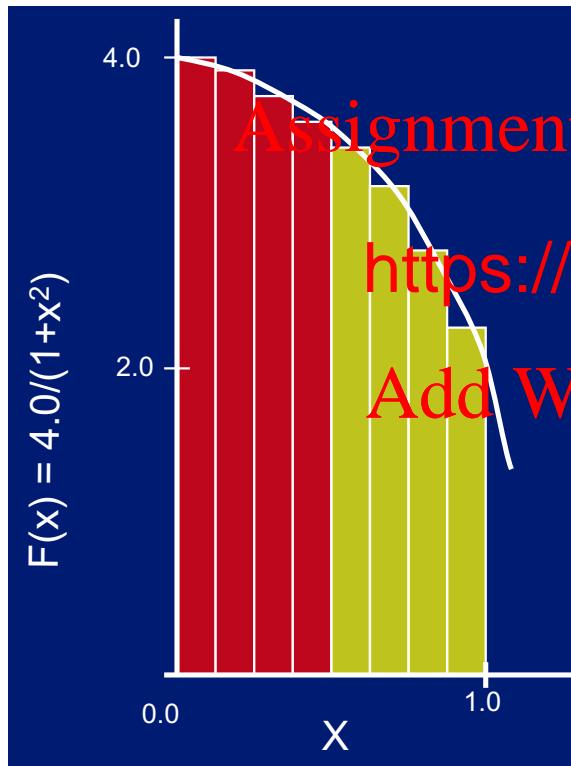
OpenMP “parallel” region

```
#pragma omp parallel
{
    Nthrds = o
    id = 0
    istart = id *
    iend = (id+
    if (id == Nthrds-1)iend = N;
    for(i=istart;i<iend;i++) { a[i] = a[i] + b[i];}
}
```

OpenMP “parallel” region and a worksharing “**for**” construct

```
#pragma omp parallel
#pragma omp for
for(i=0;i<N;i++) { a[i] = a[i] + b[i];}
```

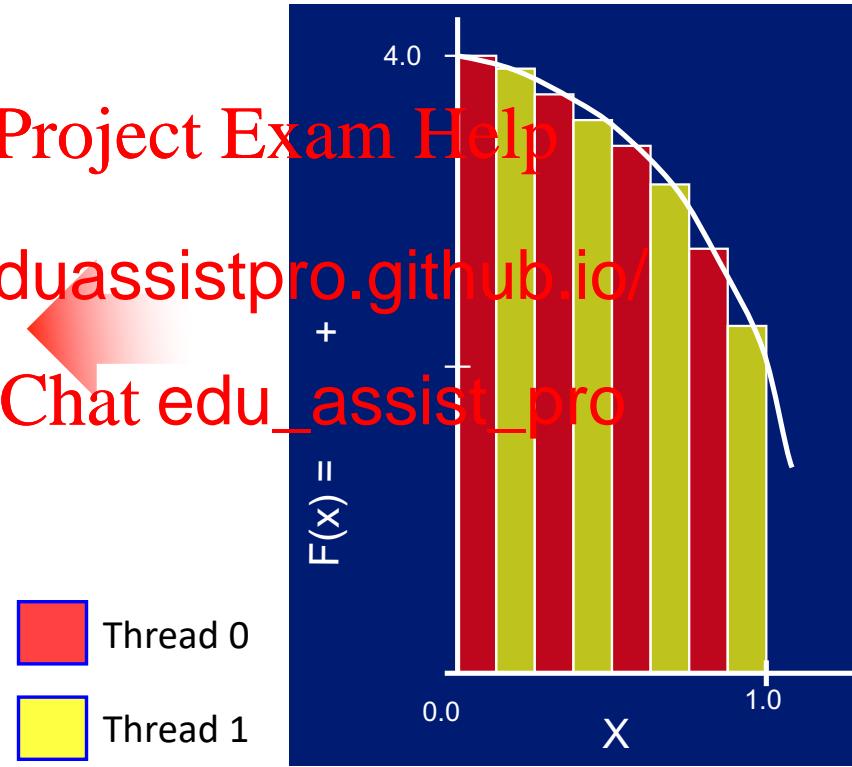
Other Concurrency Opportunities



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Example 8: Parallel For and Reduction

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double p
    step = 1.0/(dou
    #pragma omp par
    {
        #pragma omp for reduction(+:sum)
        for (i=0;i<num_steps;i++)
            x = (i+0.5)*step;
            sum = sum + 4.0/(1.0+x*x);
    }
    Pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds

1.22 seconds

Reduction is more scalable than critical sections.

Loops must not have loop carry dependency.
Limitation of compiler technology.

"i" is private by default

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Combined Parallel/Worksharing Construct

- OpenMP shortcut: Put the “parallel” and the worksharing directive on the same line

Assignment Project Exam Help

```
double res[MAX]  
#pragma omp pa  
{  
    #pragma omp for  
    for (i=0;i< MAX; i++) {  
        res[i] = huge();  
    }  
}
```

<https://eduassistpro.github.io>



```
res[MAX]; int i;  
# pragma parallel for  
i< MAX; i++) {  
    res[i] = huge();  
}
```

OpenMP: Working With Loops

- Basic approach
 - Find compute intensive loops
 - Make the loop iterations independent. So they can safely execute in any order without loop-carried dependencies
 - Place the appro

<https://eduassistpro.github.io/>

```
double hoop
for (i=id, x=(i+0.5)*step;i< num
    x += hoop;
    sum = sum + 4.0/(1.0+x*x);
}
```

x depend on
i not in for



```
for (i=0;i< num_steps; i++){
    x = (i+0.5)*step;
    sum = sum + 4.0/(1.0+x*x);
}
```

OpenMP: Reductions

```
for (i=0;i< num_steps; i++){  
    x = (i+0.5)*step;  
    sum = sum + 4.0/(1.0+x*x);  
}
```

reduction (op : list)

1. A local copy of each list variable is made and initialized depending on g. 0 for “+”)

```
#pragma omp for r  
for (i=0;i< num_s  
    x = (i+0.5)*step;  
    sum = sum + 4.0/(1.0+x*x);  
}
```

Assignment Project Exam Help



<https://eduassistpro.github.io>

union the local copy are reduced into a combined with the value

Add WeChat edu_assist_pro

- Accumulating values into a single variable (sum) creates true dependence between loop iterations that can't be trivially removed
- This is a very common situation ... it is called a “reduction”.
- Support for reduction operations is included in most parallel programming environments.

OpenMP: Reduction Operators

- Many different associative operands can be used with reduction:
- Initial values are the ones that make sense mathematically.

Assignment Project Exam Help

Operator	https://eduassistpro.github.io/
+	0
*	1
-	0

Operator	Initial value
	~0
	0
^	0
&&	1
	0

Example 8: Parallel For and Reduction

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;

int main () {
    int i; double p
    step = 1.0/(dou
    #pragma omp par
    for (i=0;i< num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x);
    }
    Pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds

Assignment Project Exam Help

1.22 seconds

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro
We created a parallel
out changing any code
and by adding 2 simple lines of text!

Disadvantage:

- 1) Lot's of pitfalls if you don't understand system architecture
- 2) The basic result may not be the fastest one can do

Outline

- Multicores Shared-Memory Platforms
- OpenMP – Parallel Shared-Memory Multiprocessing
 - Overview
 - Example: numeric integration
 - Shared variable
 - SPMD vs worksharing
 - **Data environment options**
 - Advanced worksharing options
 - Synchronization options
- Sequential Optimizations
 - Cost Model
 - Cost Reduction

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat **edu_assist_pro**

Data Environment in OpenMP

- Shared Memory programming model:
 - Most variables are shared by default

Assignment Project Exam Help

- Global variables are **SHARED** among threads
 - File scope v
 - Dynamically <https://eduassistpro.github.io/> (ew)
- But not everything is shared...
 - Functions called from parallel regions are **PRIVATE**
 - Automatic variables within a statement block are **PRIVATE**.

Add WeChat edu_assist_pro

Example: Data Sharing

```
double A[10];
int main() {
    int index[10];
    #pragma omp parallel
        work(index)
    printf("%d\n", ind
    }
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

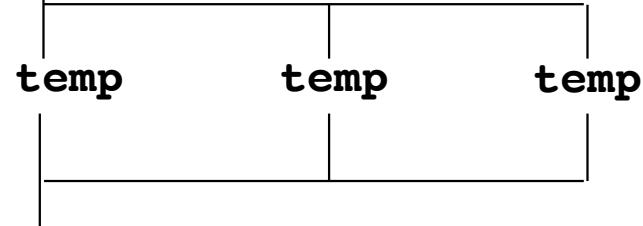
A, index and count are shared by all threads.

temp is local to each thread

```
extern double A[10];
void work(int *index) {
    double temp[10];
    static int count;
```

Add WeChat edu_assist_pro

A, index, count



Changing Storage Attributes

- One can selectively change storage attributes for constructs using the following clauses:

- SHARED
- PRIVATE
- FIRSTPRIVATE

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- The final value of a private inside transmitted to the shared variable

- LASTPRIVATE

Add WeChat edu_assist_pro

p can be
loop with:

typically use critical / atom(...)

Example 9: firstprivate example

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;
#define NUM_THREADS 2
int main ()
{
    int i; double pi, x, local_sum = 0.0; int nthrds;
    step = 1.0/(dou
    omp_set_num_thr
    #pragma omp parsum
    {
        int i, id = omp_get_thread_num()
        int nthrds = omp_get_num_threads
        double hoop = nthrds*step;
        for (i=id, x=(i+0.5)*step; i< num_steps; i=i+nthrds) {
            x += hoop;
            local_sum += 4.0/(1.0+x*x);
        }
        #pragma omp critical
        pi += local_sum * step;
    }
    printf("pi = %f \n", pi);
    return 0;
}
```

Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

Outline

- Multicores Shared-Memory Platforms
- OpenMP – Parallel Shared-Memory Multiprocessing
 - Overview
 - Example: numerics
 - Shared variable
 - SPMD vs worksharing
 - Data environment options
 - **Advanced worksharing options**
 - Synchronization options
- Sequential Optimizations
 - Cost Model
 - Cost Reduction

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The Scheduling Cause

- The schedule clause affects how loop iterations are mapped onto threads
 - schedule(static,[chunk])
 - Deal-out blocks of iterations of size “chunk” to each thread.
 - schedule(dynamic)
 - Each thread handles iterations sequentially until all iterations have been handled.
 - schedule(guided,[chunk])
 - Threads dynamically grab blocks of size “chunk”. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds.
 - schedule(runtime)
 - Schedule and chunk size taken from the OMP_SCHEDULE environment variable (or the runtime library)

The Scheduling Cause

Schedule Clause	When To Use	Help
STATIC	<p>https://eduassistpro.github.io/</p> <p>programmer</p>	Least work at runtime : scheduling done at compile-time
DYNAMIC	<p>Add WeChat edu_assist_pro</p> <p>Unpredictable, variable work per iteration</p>	Most work at runtime : complex scheduling logic used at run-time
GUIDED	<p>Special case of dynamic to reduce scheduling overhead</p>	

Example 9: Scheduling

```
#include <omp.h>
#include <stdio.h>
static long num_steps = 100000000; double step;

int main () {
    int i; double p
    step = 1.0/(dou
#pragma omp par# pragma omp parallel, schedule(static,
10000000)
    for (i=0;i< num_steps; i++){
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = sum * step;
    printf("pi = %f \n", pi);
    return 0;
}
```

Sequential: 2.12 seconds

1.31 seconds

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro
^{10 chunks}

Outline

- Multicores Shared-Memory Platforms
- OpenMP – Parallel Shared-Memory Multiprocessing
 - Overview
 - Example: numeric integration
 - Shared variable
 - SPMD vs worksharing
 - Data environment options
 - Advanced worksharing options
 - **Synchronization options**
- Sequential Optimizations
 - Cost Model
 - Cost Reduction

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Synchronization: ordered

- The **ordered** region executes in the sequential order
- Important for some scientific code and optimization code

• Order of reduction may cause rounding differences

<https://eduassistpro.github.io/>

```
#pragma omp parallel for reduction(plus: res)
for (l=0;l<N;l++){
    tmp = NEAT_STUFF(l);
    #pragma ordered
    res += consum(tmp);
}
```

Synchronization Barrier

- **Barrier:** Each thread waits until all threads arrive

Assignment Project Exam Help

```
#pragma omp parallel shared (A, B, C) private(id)
{
    id=omp_get_id();
    A[id] = big_c
    #pragma omp barrier
    #pragma omp for
    for(i=0;i<N;i++){C[i]=big_calc3(i,A);}
    #pragma omp for nowait
    for(i=0;i<N;i++){ B[i]=big_calc2(C, i); }
    A[id] = big_calc4(id);
}
```

it barrier at the end of
forksharing construct

it barrier due to
nowait

implicit barrier at the end of
a parallel region

“Single” Worksharing Construct

- The **single** construct denotes a block of code that is executed by only one thread (not necessarily the master thread).
- A barrier is implied at the end of the single block
 - can remove the barrier with a nowait clause

<https://eduassistpro.github.io/>

```
#pragma omp parallel
{
    do_many_things();
#pragma omp single
    { exchange_boundaries(); }
    do_many_other_things();
}
```

Add WeChat

```
parallel
edu_assist_pro
do_many_things();
#pragma omp single nowait
{ exchange_boundaries(); }
do_many_other_things();
}
```

Nested Parallelism

- **Q:** Is nested parallel possible with OpenMP?
- **A:** Yes. But be sure to understand why you want to use it.

```
#include <omp.h>
#include <stdio.h>
void report_num_threads()
{ printf("I have %d threads\n"); }
int main()
{
    omp_set_dynamic(0);
    #pragma omp parallel num_threads(2)
    {
        report_num_threads();
        #pragma omp parallel num_threads(2)
        {
            report_num_threads();
            #pragma omp parallel num_threads(2)
            report_num_threads();
        }
    }
    return(0);
}
```

Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

<http://download.oracle.com/docs/cd/E19205-01/819-5270/aewbc/index.html>

Nested Parallelism

- Compiling and running this program with nested parallelism enabled produces the following (sorted) output:

Assignment Project Exam Help

```
$ export OMP_NESTED=TRUE
$ ./experimentN
L 1: # threads in team 1
L 2: # threads in team 2
L 2: # threads in team 2
L 3: # threads in team 2
```

NESTED=FALSE

```
n team 1
n team 1
n team 1
n team 1
```

OpenMP Environment Variables

- OMP_SCHEDULE=algorithm

- dynamic[, n]
 - guided[, n]
 - Runtime
 - static[, n]

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- OMP_NUM_THREADS=num

Add WeChat edu_assist_pro

- OMP_NESTED=TRUE|FALSE

- OMP_DYNAMIC=TRUE|FALSE

Outline

- Multicores Shared-Memory Platforms
- OpenMP – Parallel Shared-Memory Multiprocessing
 - Overview
 - Example: numerics
 - Shared variable
 - SPMD vs worksharing
 - Data environment options
 - Advanced worksharing options
 - Synchronization options
- Sequential Optimizations
 - Cost Model
 - Cost Reduction

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Optimizing the Base Case

- Base case is usually the guts of the algorithm that computes on data in caches

Assignment Project Exam Help

- Should fit in regis

<https://eduassistpro.github.io/>

- Must use SIMD

Add WeChat edu_assist_pro

Cost Model

- **Algebraic model**

- +, -, /, *

- sin, cos, etc.

```
int func(int i) {  
    return i+2;  
}
```

- **Realistic model**

- +, -, /, *

- std libs (sin, cos, etc)

- array[i]

- A.sum, A.prod

- (double) x

- func(...)

- i < n, compute

- i < n, test and branch

- {} unconditional branch

prod;

<https://eduassistpro.github.io/>

Add WeChat [edu_assist_pro](https://eduassistpro.github.io/)

```
A a array[] {  
    A A .0},  
    for(int i=0; i<n; ++i) {  
        A.sum = A.sum +  
        (double)func(array[i]);  
        A.prod = A.prod *  
        (double)func(array[i]);  
    }  
    return res;  
}
```

Cost Model (cont'd)

- Algebraic model

cost = **3 n**

```
int func(int i) {  
    return i+2;  
}
```

- Realistic model

+ and *

2 int -> double

4 a.x

++i

i < n conditional jump

2 func() calls

cost = **19 n**

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
A acc(in  
A res  
for(in  
    A.sum = A.sum + (double)func(i);  
    A.prod = A.prod * (double)func(i);  
}  
return res;  
}
```

6x discrepancy in cost

Using the Cost Model to Optimize

```
int func(int i) { return i+2; }
```

```
struct A { double sum, prod; };
```

Assignment Project Exam Help

```
A acc
```

```
A r
```

<https://eduassistpro.github.io/>

Cost = 19n

```
for(int i=0; i<n; ++i)
    A.sum = A.sum + (dou
    A.prod = A.prod * (d
}
return res;

}
```

Add WeChat edu_assist_pro

Using the Cost Model to Optimize I

```
int func(int i) { return i+2; }
```

```
struct A { double sum, prod; };
```

Assignment Project Exam Help

```
A acc
```

```
dou https://eduassistpro.github.io/  
for(int i=0; i<n; ++i)  
    f double (i+1)  
    sum = sum + f;  
    prod = prod * f;  
}  
A res = {sum, prod}  
return res;  
}
```

Cost = $6n$

Add WeChat edu_assist_pro

Cost Reduction = closing down the gap

Using the Cost Model to Optimize II

```
int func(int i) { return i+2; }
```

```
struct A { double sum, prod; };
```

Assignment Project Exam Help

```
A acc
```

```
dou https://eduassistpro.github.io/  
for(int i=2; i<n+2; ++)  
    f = (double)(i)  
    sum = sum + f;  
    prod = prod * f;  
}  
A res = {sum, prod};  
return res;  
}
```

Cost = $5n$

Add WeChat edu_assist_pro

Optimizations enable new optimizations

Using the Cost Model to Optimize III

```
int func(int i) { return i+2; }

struct A { double sum, prod; };
A acc(int n, int array[]) {
    dou #pr https://eduassistpro.github.io/
    for(int i=2; i<n+2; ++
        f = (double)(i)
        sum = sum + f;
        prod = prod * f;
    }
    A res = {sum, prod};
    return res;
}
```

Assignment Project Exam Help

#pr <https://eduassistpro.github.io/>

Cost = $3.5n$

Add WeChat edu_assist_pro

Final result: $3.5n$ vs $19n$

Outline

- Multicores Shared-Memory Platforms
- OpenMP – Parallel Shared-Memory Multiprocessing
 - Overview
 - Example: numeric integration
 - Shared variable
 - SPMD vs worksharing
 - Data environment options
 - Advanced worksharing options
 - Synchronization options
- Sequential Optimizations
 - Cost Model
 - **Cost Reduction**

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Cost Reduction: A Compiler Problem

- **Solution:** “Human Compiler”
- Compilers often fail on complex codes (many assumptions are violated)

Assignment Project Exam Help

- Optimizations
 - Strength reduction
 - Function inlining
 - Loop unrolling
 - Common subexpression elimination
 - Load/store elimination
 - Table lookups
 - Branch elimination

Add WeChat edu_assist_pro

Strength Reduction

```
for i = 1..n
```

```
    sum = sum + a[i] / c;
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

```
double c_inv = 1/c;  
for i = 1..n  
    sum = sum + a[i] * c_inv;
```

division is
more expensive

than multiplication

Expensive operations:
/, %, sin, cos, log

Function Inlining

```
double f(a) { return a/c; }
```

Assignment Project Exam Help
for i = 1..n

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
for i = 1..n  
    sum = sum + a[i] / c;
```

Now strength reduction can apply

Loop Unrolling

```
for i = 1..n
```

```
    sum = sum + a[i]
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

 *sequential locality*
for i = 1:4:n
 sum = sum + a[i] + a[i+1] + a[i+2] + a[i+3]

```
// remaining iterations
for i = i:n
    sum = sum + a[i]
```

Often enables other optimizations

Common Sub-expression Elimination

```
for j = 1..m
```

```
for i = 1..n
```

```
sum = sum + cos(j*PI/180)*a[i];
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

for j = 1..m

```
double c = cos(j*PI/180);
```

```
for i = 1..n
```

```
sum = sum + c*a[i];
```

Add WeChat edu_assist_pro

Table Lookups

```
for j = 1..m
```

```
    for i = 1..n
```

```
        sum = sum + cos(j*PI/180)*a[i];
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
for j = 1..m
    double c = COS_TAB[j];
    for i = 1..n
        sum = sum + c*a[i];
```

Load/Store Elimination

Loop Merging

Assignment Project Exam Help

```
for i = 1..n  
    sum = sum + f(a[i]);
```

```
for i
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
for i = 1..n  
    sum = sum + f(a[i]);  
    prod = prod * f(a[i]);
```

One of the most important optimizations!
Almost always enables other optimizations.

Branch Elimination

```
for i = 1..n  
    if(i != except)  
        sum = sum + a[i];
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

```
for i = 1..except-1  
    sum = sum + a[i];  
for i = except+1..n  
    sum = sum + a[i];
```

~~sum = 0.0;~~

What we have discussed in the last 2 lectures

Fast Platforms

Good Techniques

Assignment Project Exam Help

- Multicore platforms
- Manycore
- Cloud platforms

structures

<https://eduassistpro.github.io/>
Architecture

- OpenMP Abstractions:
 - Help establish a mental model for the programmers
 - Make it easier to write parallel code
 - Performance depends on deep understanding of the implementation platform